

THE C++ PROGRAMMING LANGUAGE – BJARNE STROUSTRUP

(ADVICE LIST)

INTRODUCTORY MATERIAL

NOTES TO THE READER: PROGRAMMING IN C++

- Represent ideas directly in code.
- Represent relationships among ideas directly in code (e.g., hierarchical, parametric, and ownership relationships).
- Represent independent ideas independently in code.
- Keep simple things simple (without making complex things impossible).
- Prefer statically type-checked solutions (when applicable).
- Keep information local (e.g., avoid global variables, minimize the use of pointers).
- Don't over abstract (i.e., don't generalize, introduce class hierarchies, or parameterize beyond obvious needs and experience)

NOTES TO THE READER: SUGGESTIONS FOR (OLD & NEW) C++ PROGRAMMERS:

- Use constructors to establish invariants (§2.4.3.2, §13.4, §17.2.1).
- Use constructor/destructor pairs to simplify resource management (RAII; §5.2, §13.3).
- Avoid “naked” new and delete (§3.2.1.2, §11.2.1).
- Use containers and algorithms rather than built-in arrays and ad hoc code (§4.4, §4.5, §7.4, Chapter 32).
- Prefer standard-library facilities to locally developed code (§1.2.4).
- Use exceptions, rather than error codes, to report errors that cannot be handled locally (§2.4.3, §13.1).
- Use move semantics to avoid copying large objects (§3.3.2, §17.5.2).
- Use `unique_ptr` to reference objects of polymorphic type (§5.2.1).
- Use `shared_ptr` to reference shared objects, that is, objects without a single owner that is responsible for their destruction (§5.2.1).
- Use templates to maintain static type safety (eliminate casts) and avoid unnecessary use of class hierarchies (§27.2).

NOTES TO THE READER: SUGGESTIONS FOR C PROGRAMMERS: The better one knows C, the harder it seems to be to avoid writing C++ in C style, thereby losing many of the potential benefits of C++:

- Don't think of C++ as C with a few features added. C++ can be used that way, but only suboptimally. To get really major advantages from C++ as compared to C, you need to apply different design and implementation styles.
- Don't write C in C++; that is often seriously suboptimal for both maintenance and performance.
- Use the C++ standard library as a teacher of new techniques and programming styles. Note the difference from the C standard library (e.g., `=` rather than `strcpy()` for copying and `==` rather than `strcmp()` for comparing).
- Macro substitution is almost never necessary in C++. Use `const`, `constexpr`, `enum` or `enum class` to define manifest constants, `inline` to avoid function-calling overhead, `template s` (§to specify families of functions and types, and `namespace s` to avoid name clashes.
- Don't declare a variable before you need it, and initialize it immediately. A declaration can occur anywhere a statement can, in for-statement initializers (§9.5), and in conditions.
- Don't use `malloc()`. The `new` operator does the same job better, and instead of `realloc()`, try a `vector`. Don't just replace `malloc()` and `free()` with “naked” `new` and `delete`.
- Avoid `void*`, unions, and casts, except deep within the implementation of some function or class. Their use limits the support you can get from the type system and can harm performance. In most cases, a cast is an indication of a design error. If you must use an explicit type conversion, try using one of the named casts (e.g., `static_cast;`) for a more precise statement of what you are trying to do.
- Minimize the use of arrays and C-style strings. C++ standard-library `string s`, `array s`, and `vector s` can often be used to write simpler and more maintainable code compared to the traditional C style. In general, try not to build yourself what has already been provided by the standard library.
- Avoid pointer arithmetic except in very specialized code (such as a memory manager) and for simple array traversal (e.g., `++p`).
- Do not assume that something laboriously written in C style (avoiding C++ features such as classes, templates, and exceptions) is more efficient than a shorter alternative (e.g., using standard-library facilities). Often (but of course not always), the opposite is true.

NOTES TO THE READER: SUGGESTIONS FOR JAVA PROGRAMMERS (MOST OF THEM APPLIABLE FOR C# AS WELL)

- Don't simply mimic Java style in C++; that is often seriously suboptimal for both maintainability and performance
- Use the C++ abstraction mechanisms (e.g. class and templates): don't fall back to a C style of programming out of a false feeling of familiarity
- Use the C++ standard library as a teacher of new techniques and programming styles.
- Don't immediately invent a unique base for all of your classes (an Object class). Typically, you can do better without it for many/most classes.
- Minimize the use of reference and pointer variables: use local and member variables.
- Remember: a variable is never implicitly a reference.
- Think of pointers as C++'s equivalent to Java references (C++ references are more limited; there is no reseating of C++ references).
- A function is not virtual by default. Not every class is meat for inheritance.
- Use abstract classes as interfaces to class hierarchies; avoid "brittle base classes", that is, base classes with data members.
- Use scoped resource management ("Resource Acquisition Is Initialization"; RAII) whenever possible.
- Use a constructor to establish a class invariant (and throw an exception if it can't)
- If a cleanup action is needed when an object is deleted (e.g., goes out of scope), use a destructor for that. Don't imitate finally (doing so is more ad hoc in the longer run far more work than relying on destructors).
- Avoid "naked new and delete; instead, use containers (e.g. vector, string, and map) and handle classes (e.g. lock and unique_ptr).
- Use freestanding functions (nonmembers functions) to minimize coupling (e.g., see the standard algorithms), and use namespaces to limit the scope of freestanding functions.
- Don't use exception specifications (except no except)
- A C++ nested class does not have access to an object of the enclosing class.
- C++ offers only the most minimal run-time reflection: dynamic_cast and typeid. Rely more on compile-time facilities (e.g., compile-time polymorphism)

NOTES TO THE READER

- Represent ideas (concepts) directly in code, for example, as a function, a class, or an enumeration; §1.2.
- Aim for your code to be both elegant and efficient; §1.2.
- Don't over abstract; §1.2.
- Focus design on the provision of elegant and efficient abstractions, possibly presented as libraries; §1.2.
- Represent relationships among ideas directly in code, for example, through parameterization or a class hierarchy; §1.2.1.
- Represent independent ideas separately in code, for example, avoid mutual dependencies among classes; §1.2.1.
- C++ is not just object-oriented; §1.2.1.
- C++ is not just for generic programming; §1.2.1.
- Prefer solutions that can be statically checked; §1.2.1.
- Make resources explicit (represent them as class objects); §1.2.1, §1.4.2.1.
- Express simple ideas simply; §1.2.1.
- Use libraries, especially the standard library, rather than trying to build everything from scratch; §1.2.1.
- Use a type-rich style of programming; §1.2.2.
- Low-level code is not necessarily efficient; don't avoid classes, templates, and standard-library components out of fear of performance problems; §1.2.4, §1.3.3.
- If data has an invariant, encapsulate it; §1.3.2.
- C++ is not just C with a few extensions; §1.3.3. In general: To write a good program takes intelligence, taste, and patience. You are not going to get it right the first time. Experiment!

A TOUR OF C++: THE BASICS

- Don't panic! All will become clear in time; §2.1.
- You don't have to know every detail of C++ to write good programs; §1.3.1.
- Focus on programming techniques, not on language features; §2.1.

A TOUR OF C++: ABSTRACTION MECHANISMS

- Express ideas directly in code; §3.2.
- Define classes to represent application concepts directly in code; §3.2.
- Use concrete classes to represent simple concepts and performance-critical components; §3.2.1.
- Avoid “naked” new and delete operations; §3.2.1.2.
- Use resource handles and RAII to manage resources; §3.2.1.2.
- Use abstract classes as interfaces when complete separation of interface and implementation is needed; §3.2.2.
- Use class hierarchies to represent concepts with inherent hierarchical structure; §3.2.4.
- When designing a class hierarchy, distinguish between implementation inheritance and interface inheritance; §3.2.4.
- Control construction, copy, move, and destruction of objects; §3.3.
- Return containers by value (relying on move for efficiency); §3.3.2.
- Provide strong resource safety; that is, never leak anything that you think of as a resource; §3.3.3.
- Use containers, defined as resource handle templates, to hold collections of values of the same type; §3.4.1.
- Use function templates to represent general algorithms; §3.4.2.
- Use function objects, including lambdas, to represent policies and actions; §3.4.3.
- Use type and template aliases to provide a uniform notation for types that may vary among similar types or among implementations; §3.4.5

A TOUR OF C++: CONTAINERS AND ALGORITHMS

- Don’t reinvent the wheel; use libraries; §4.1.
- When you have a choice, prefer the standard library over other libraries; §4.1.
- Do not think that the standard library is ideal for everything; §4.1.
- Remember to #include the headers for the facilities you use; §4.1.2.
- Remember that standard-library facilities are defined in namespace std; §4.1.2.
- Prefer strings over C-style strings (uchar*; §2.2.5); §4.2, §4.3.2.
- [7]iostreams are type sensitive, type-safe, and extensible; §4.3.
- Prefer vector<T>, map<K,T>, and unordered_map<K,T> over T[]; §4.4.
- Know your standard containers and their tradeoffs; §4.4.
- Use vector as your default container; §4.4.1.
- Prefer compact data structures; §4.4.1.1.
- If in doubt, use a range-checked vector (such as Vec); §4.4.1.2.
- Use push_back() or back_inserter() to add elements to a container; §4.4.1, §4.5.
- Use push_back() on a vector rather than realloc() on an array; §4.5.
- Catch common exceptions in main(); §4.4.1.2.
- Know your standard algorithms and prefer them over handwritten loops; §4.5.5.
- If iterator use gets tedious, define container algorithms; §4.5.6

A TOUR OF C++: CONCURRENCY AND UTILITIES

- Use resource handles to manage resources (RAII); §5.2.
- Use unique_ptr to refer to objects of polymorphic type; §5.2.1.
- Use shared_ptr to refer to shared objects; §5.2.1.
- Use type-safe mechanisms for concurrency; §5.3.
- Minimize the use of shared data; §5.3.4.
- Don’t choose shared data for communication because of “efficiency” without thought and preferably not without measurement; §5.3.4.
- Think in terms of concurrent tasks, rather than threads; §5.3.5.
- A library doesn’t have to be large or complicated to be useful; §5.4.
- Time your programs before making claims about efficiency; §5.4.1.
- You can write code to explicitly depend on properties of types; §5.4.2.
- Use regular expressions for simple pattern matching; §5.5.
- Don’t try to do serious numeric computation using only the language; use libraries; §5.6.
- Properties of numeric types are accessible through numeric_limits; §5.6.5

BASIC FACILITIES

TYPES AND DECLARATIONS

- For the final word on language definition issues, see the ISO C++ standard; §6.1.
- Avoid unspecified and undefined behavior; §6.1.

- Isolate code that must depend on implementation-defined behavior; §6.1.
- Avoid unnecessary assumptions about the numeric value of characters; §6.2.3.2, §10.5.2.1.
- Remember that an integer starting with a 0 is octal; §6.2.4.1.
- Avoid “magic constants”; §6.2.4.1.
- Avoid unnecessary assumptions about the size of integers; §6.2.8.
- Avoid unnecessary assumptions about the range and precision of floating-point types; §6.2.8.
- Prefer plain char over signed char and unsigned char; §6.2.3.1.
- Beware of conversions between signed and unsigned types; §6.2.3.1.
- Declare one name (only) per declaration; §6.3.2.
- Keep common and local names short, and keep uncommon and nonlocal names longer; §6.3.3.
- Avoid similar-looking names; §6.3.3.
- Name an object to reflect its meaning rather than its type; §6.3.3.
- Maintain a consistent naming style; §6.3.3.
- Avoid ALL_CAPS names; §6.3.3.
- Keep scopes small; §6.3.4.
- Don’t use the same name in both a scope and an enclosing scope; §6.3.4.
- Prefer the {}-initializer syntax for declarations with a named type; §6.3.5.
- Prefer the = syntax for the initialization in declarations using auto; §6.3.5.
- Avoid uninitialized variables; §6.3.5.1.
- Use an alias to define a meaningful name for a built-in type in cases in which the built-in type used to represent a value might change; §6.5.
- Use an alias to define synonyms for types; use enumerations and classes to define new types; §6.5

POINTERS, ARRAYS AND REFERENCES

- Keep use of pointers simple and straightforward; §7.4.1.
- Avoid nontrivial pointer arithmetic; §7.4.
- Take care not to write beyond the bounds of an array; §7.4.1.
- Avoid multidimensional arrays; define suitable containers instead; §7.4.2.
- Use nullptr rather than 0 or NULL; §7.2.2.
- Use containers (e.g., vector, array, and valarray) rather than built-in (C-style) arrays; §7.4.1.
- Use string rather than zero-terminated arrays of char; §7.4.
- Use raw strings for string literals with complicated uses of backslash; §7.3.2.1.
- Prefer const reference arguments to plain reference arguments; §7.7.3.
- Use rvalue references (only) for forwarding and move semantics; §7.7.2.
- Keep pointers that represent ownership inside handle classes; §7.6.
- Avoid void* except in low-level code; §7.2.1.
- Use const pointers and const references to express immutability in interfaces; §7.5.
- Prefer references to pointers as arguments, except where “no object” is a reasonable option; §7.7.4

STRUCTURES, UNIONS, AND ENUMERATIONS

- When compactness of data is important, lay out structure data members with larger members before smaller ones; §8.2.1.
- Use bit-fields to represent hardware-imposed data layouts; §8.2.7.
- Don’t naively try to optimize memory consumption by packing several values into a single byte; §8.2.7.
- Use unions to save space (represent alternatives) and never for type conversion; §8.3.
- Use enumerations to represent sets of named constants; §8.4.
- Prefer class enums over “plain”enums to minimize surprises; §8.4.
- Define operations on enumerations for safe and simple use; §8.4.1.

STATEMENTS

- Don’t declare a variable until you have a value to initialize it with; §9.3, §9.4.3, §9.5.2.
- Prefer a switch-statement to an if-statement when there is a choice; §9.4.2.
- Prefer a range-for-statement to a for-statement when there is a choice; §9.5.1.
- Prefer a for-statement to a while-statement when there is an obvious loop variable; §9.5.2.
- Prefer a while-statement to a for-statement when there is no obvious loop variable; §9.5.3.
- Avoid do-statements; §9.5.
- Avoid goto; §9.6.
- Keep comments crisp; §9.7.

- Don't say in comments what can be clearly stated in code; §9.7.
- State intent in comments; §9.7.
- Maintain a consistent indentation style; §9.7

EXPRESSIONS

- Prefer the standard library to other libraries and to “handcrafted code”; §10.2.8.
- Use character-level input only when you have to; §10.2.3.
- When reading, always consider ill-formed input; §10.2.3.
- Prefer suitable abstractions (classes, algorithms, etc.) to direct use of language features (e.g., ints, statements); §10.2.8.
- Avoid complicated expressions; §10.3.3.
- If in doubt about operator precedence, parenthesize; §10.3.3.
- Avoid expressions with undefined order of evaluation; §10.3.2. Avoid narrowing conversions; §10.5.2.
- Define symbolic constants to avoid “magic constants”; §10.4.1.
- Avoid narrowing conversions; §10.5.2

SELECT OPERATIONS

- Prefer prefix++ over suffix++; §11.1.4.
- Use resource handles to avoid leaks, premature deletion, and double deletion; §11.2.1.
- Don't put objects on the free store if you don't have to; prefer scoped variables; §11.2.1.
- Avoid “naked new” and “naked delete”; §11.2.1.
- Use RAI; §11.2.1.
- Prefer a named function object to a lambda if the operation requires comments; §11.4.2.
- Prefer a named function object to a lambda if the operation is generally useful; §11.4.2.
- Keep lambdas short; §11.4.2.
- For maintainability and correctness, be careful about capture by reference; §11.4.3.1.
- Let the compiler deduce the return type of a lambda; §11.4.4.
- Use the T{e} notation for construction; §11.5.1.
- Avoid explicit type conversion (casts); §11.5.
- When explicit type conversion is necessary, prefer a named cast; §11.5.
- Consider using a run-time checked cast, such as narrow_cast<>(), for conversion between numeric types; §11.5.

FUNCTIONS

- “Package” meaningful operations as carefully named functions; §12.1.
- A function should perform a single logical operation; §12.1.
- Keep functions short; §12.1.
- Don't return pointers or references to local variables; §12.1.4.
- If a function may have to be evaluated at compile time, declare it constexpr; §12.1.6.
- If a function cannot return, mark it [[noreturn]]; §12.1.7.
- Use pass-by-value for small objects; §12.2.1.
- Use pass-by-const-reference to pass large values that you don't need to modify; §12.2.1.
- Return a result as a return value rather than modifying an object through an argument; §12.2.1.
- Use rvalue references to implement move and forwarding; §12.2.1.
- Pass a pointer if “no object” is a valid alternative (and represent “no object” by nullptr); §12.2.1.
- Use pass-by-non-const-reference only if you have to; §12.2.1.
- Use const extensively and consistently; §12.2.1.
- Assume that a char* or a const char* argument points to a C-style string; §12.2.2.
- Avoid passing arrays as pointers; §12.2.2.
- Pass a homogeneous list of unknown length as an initializer_list<T> (or as some other container); §12.2.3.
- Avoid unspecified numbers of arguments (...); §12.2.4.
- Use overloading when functions perform conceptually the same task on different types; §12.3.
- When overloading on integers, provide functions to eliminate common ambiguities; §12.3.5.
- Specify preconditions and postconditions for your functions; §12.4.
- Prefer function objects (including lambdas) and virtual functions to pointers to functions; §12.5.
- Avoid macros; §12.6.

- If you must use macros, use ugly names with lots of capital letters; §12.6

EXCEPTION HANDLING

- Develop an error-handling strategy early in a design; §13.1.
- Throw an exception to indicate that you cannot perform an assigned task; §13.1.1.
- Use exceptions for error handling; §13.1.4.2.
- Use purpose-designed user-defined types as exceptions (not built-in types); §13.1.1.
- If you for some reason cannot use exceptions, mimic them; §13.1.5.
- Use hierarchical error handling; §13.1.6.
- Keep the individual parts of error handling simple; §13.1.6.
- Don't try to catch every exception in every function; §13.1.6.
- Always provide the basic guarantee; §13.2, §13.6.
- Provide the strong guarantee unless there is a reason not to; §13.2, §13.6.
- Let a constructor establish an invariant, and throw if it cannot; §13.2.
- Release locally owned resources before throwing an exception; §13.2.
- Be sure that every resource acquired in a constructor is released when throwing an exception in that constructor; §13.3.
- Don't use exceptions where more local control structures will suffice; §13.1.4.
- Use the "Resource Acquisition Is Initialization" technique to manage resources; §13.3.
- Minimize the use of try-blocks; §13.3.
- Not every program needs to be exception-safe; §13.1.
- Use "Resource Acquisition Is Initialization" and exception handlers to maintain invariants; §13.5.2.2.
- Prefer proper resource handles to the less structured finally; §13.3.1.
- Design your error-handling strategy around invariants; §13.4.
- What can be checked at compile time is usually best checked at compile time (using static_assert); §13.4.
- Design your error-handling strategy to allow for different levels of checking/enforcement; §13.4.
- If your function may not throw, declare it noexcept; §13.5.1.1
- Don't use exception specification; §13.5.1.3.
- Catch exceptions that may be part of a hierarchy by reference; §13.5.2.
- Don't assume that every exception is derived from class exception; §13.5.2.2.
- Have main() catch and report all exceptions; §13.5.2.2, §13.5.2.4.
- Don't destroy information before you have its replacement ready; §13.6.
- Leave operands in valid states before throwing an exception from an assignment; §13.2.
- Never let an exception escape from a destructor; §13.2.
- Keep ordinary code and error-handling code separate; §13.1.1, §13.1.4.2.
- Beware of memory leaks caused by memory allocated by new not being released in case of an exception; §13.3.
- Assume that every exception that can be thrown by a function will be thrown; §13.2.
- A library shouldn't unilaterally terminate a program. Instead, throw an exception and let a caller decide; §13.4.
- A library shouldn't produce diagnostic output aimed at an end user. Instead, throw an exception and let a caller decide; §13.1.3

NAMESPACES

- Use namespaces to express logical structure; §14.3.1.
- Place every nonlocal name, except main(), in some namespace; §14.3.1.
- Design a namespace so that you can conveniently use it without accidentally gaining access to unrelated namespaces; §14.3.3.
- Avoid very short names for namespaces; §14.4.2.
- If necessary, use namespace aliases to abbreviate long namespace names; §14.4.2.
- Avoid placing heavy notational burdens on users of your namespaces; §14.2.2, §14.2.3.
- Use separate namespaces for interfaces and implementations; §14.3.3.
- Use the Namespace::member notation when defining namespace members; §14.4.
- Use inline namespaces to support versioning; §14.4.6.
- Use using-directives for transition, for foundational libraries (such as std), or within a local scope; §14.4.9.
- Don't put a using-directive in a header file; §14.2.3

SOURCE FILES AND PROGRAMS

- Use header files to represent interfaces and to emphasize logical structure; §15.1, §15.3.2.
- #include a header in the source file that implements its functions; §15.3.1.

- Don't define global entities with the same name and similar-but-different meanings in different translation units; §15.2.
- Avoid non-inline function definitions in headers; §15.2.2.
- Use #include only at global scope and in namespaces; §15.2.2.
- #include only complete declarations; §15.2.2.
- Use include guards; §15.3.3.
- #include C headers in namespaces to avoid global names; §14.4.9, §15.2.4.
- Make headers self-contained; §15.2.3.
- Distinguish between users' interfaces and implementers' interfaces; §15.3.2.
- Distinguish between average users' interfaces and expert users' interfaces; §15.3.2.
- Avoid nonlocal objects that require run-time initialization in code intended for use as part of non-C++ programs; §15.4.1

ABSTRACTION MECHANISMS

CLASSES

- Represent concepts as classes; §16.1.
- Separate the interface of a class from its implementation; §16.1. Use public data (structs) only when it really is just data and no invariant is meaningful for the data members; §16.2.4. Define a constructor to handle initialization of objects; §16.2.5.
- By default declare single-argument constructors explicit; §16.2.6. Declare a member function that does not modify the state of its object const; §16.2.9.
- A concrete type is the simplest kind of class. Where applicable, prefer a concrete type over more complicated classes and over plain data structures; §16.3.
- Make a function a member only if it needs direct access to the representation of a class; §16.3.2.
- Use a namespace to make the association between a class and its helper functions explicit; §16.3.2.
- Make a member function that doesn't modify the value of its object a const member function; §16.2.9.1.
- Make a function that needs access to the representation of a class but needn't be called for a specific object a static member function; §16.2.12.

CONSTRUCTION, CLEANUP, COPY AND MOVE

- Design constructors, assignments, and the destructor as a matched set of operations; §17.1.
- Use a constructor to establish an invariant for a class; §17.2.1.
- If a constructor acquires a resource, its class needs a destructor to release the resource; §17.2.2.
- If a class has a virtual function, it needs a virtual destructor; §17.2.5.
- If a class does not have a constructor, it can be initialized by member wise initialization; §17.3.1.
- Prefer {}-initialization over = and () initialization; §17.3.2.
- Give a class a default constructor if and only if there is a "natural" default value; §17.3.3.
- If a class is a container, give it an initializer-list constructor; §17.3.4.
- Initialize members and bases in their order of declaration; §17.4.1.
- If a class has a reference member, it probably needs copy operations (copy constructor and copy assignment); §17.4.1.1.
- Prefer member initialization over assignment in a constructor; §17.4.1.1.
- Use in-class initializers to provide default values; §17.4.4.
- If a class is a resource handle, it probably needs copy and move operations; §17.5.
- When writing a copy constructor, be careful to copy every element that needs to be copied (beware of default initializers); §17.5.1.1.
- A copy operations should provide equivalence and independence; §17.5.1.3.
- Beware of entangled data structures; §17.5.1.3.
- Prefer move semantics and copy-on-write to shallow copy; §17.5.1.3.
- If a class is used as a base class, protect against slicing; §17.5.1.4.
- If a class needs a copy operation or a destructor, it probably needs a constructor, a destructor, a copy assignment, and a copy constructor; §17.6.
- If a class has a pointer member, it probably needs a destructor and non-default copy operations; §17.6.3.3.
- If a class is a resource handle, it needs a constructor, a destructor, and non-default copy operations; §17.6.3.3.
- If a default constructor, assignment, or destructor is appropriate, let the compiler generate it (don't rewrite it yourself); §17.6.
- Be explicit about your invariants; use constructors to establish them and assignments to maintain them; §17.6.3.2.
- Make sure that copy assignments are safe for self-assignment; §17.5.1.

- When adding a new member to a class, check to see if there are user-defined constructors that need to be updated to initialize the member; §17.5.1

OVERLOADING

- Define operators primarily to mimic conventional usage; §18.1.
- Redefine or prohibit copying if the default is not appropriate for a type; §18.2.2.
- For large operands, use const reference argument types; §18.2.4.
- For large results, use a move constructor; §18.2.4.
- Prefer member functions over nonmembers for operations that need access to the representation; §18.3.1.
- Prefer nonmember functions over members for operations that do not need access to the representation; §18.3.2.
- Use namespaces to associate helper functions with “their” class; §18.2.5.
- Use nonmember functions for symmetric operators; §18.3.2.
- Use member functions to express operators that require an lvalue as their left-hand operand; §18.3.3.1.
- Use user-defined literals to mimic conventional notation; §18.3.4.
- Provide “set() and get() functions” for a data member only if the fundamental semantics of a class require them; §18.3.5.
- Be cautious about introducing implicit conversions; §18.4.
- Avoid value-destroying (“narrowing”) conversions; §18.4.1.
- Do not define the same conversion as both a constructor and a conversion operator; §18.4.3

SPECIAL OPERATIONS

- Use operator[]() for subscripting and for selection based on a single value; §19.2.1.
- Use operator()() for call semantics, for subscripting, and for selection based on multiple values; §19.2.2.
- Use operator->() to dereference “smart pointers”; §19.2.3.
- Prefer prefix++ over suffix++; §19.2.4.
- Define the global operator new() and operator delete() only if you really have to; §19.2.5.
- Define member operator new() and member operator delete() to control allocation and de-allocation of objects of a specific class or hierarchy of classes; §19.2.5.
- Use user-defined literals to mimic conventional notation; §19.2.6.
- Place literal operators in separate namespaces to allow selective use; §19.2.6.
- For nonspecialized uses, prefer the standard string (Chapter 36) to the result of your own exercises; §19.3.
- Use a friend function if you need a nonmember function to have access to the representation of a class (e.g., to improve notation or to access the representation of two classes); §19.4.
- Prefer member functions to friend functions for granting access to the implementation of a class; §19.4.2.

DERIVED CLASSES

- Avoid type fields; §20.3.1.
- Access polymorphic objects through pointers and references; §20.3.2.
- Use abstract classes to focus design on the provision of clean interfaces; §20.4.
- Use override to make overriding explicit in large class hierarchies; §20.3.4.1.
- Use final only sparingly; §20.3.4.2.
- Use abstract classes to specify interfaces; §20.4.
- Use abstract classes to keep implementation details out of interfaces; §20.4.
- A class with a virtual function should have a virtual destructor; §20.4.
- An abstract class typically doesn’t need a constructor; §20.4.
- Prefer private members for implementation details; §20.5.
- Prefer public members for interfaces; §20.5.
- Use protected members only carefully when really needed; §20.5.1.1.
- Don’t declare data members protected; §20.5.1.1.

CLASS HIERARCHIES

- Use unique_ptr or shared_ptr to avoid forgetting to delete objects created using new; §21.2.1.
- Avoid data members in base classes intended as interfaces; §21.2.1.1.
- Use abstract classes to express interfaces; §21.2.2.
- Give an abstract class a virtual destructor to ensure proper cleanup; §21.2.2.
- Use override to make overriding explicit in large class hierarchies; §21.2.2.

- Use abstract classes to support interface inheritance; §21.2.2.
- Use base classes with data members to support implementation inheritance; §21.2.2.
- Use ordinary multiple inheritance to express a union of features; §21.3.
- Use multiple inheritance to separate implementation from interface; §21.3.
- Use a virtual base to represent something common to some, but not all, classes in a hierarchy; §21.3.5

RUN-TIME TYPE INFORMATION

- Use virtual functions to ensure that the same operation is performed independently of which interface is used for an object; §22.1.
- Use `dynamic_cast` where class hierarchy navigation is unavoidable; §22.2.
- Use `dynamic_cast` for type-safe explicit navigation of a class hierarchy; §22.2.1.
- Use `dynamic_cast` to a reference type when failure to find the required class is considered a failure; §22.2.1.1.
- Use `dynamic_cast` to a pointer type when failure to find the required class is considered a valid alternative; §22.2.1.1.
- Use double dispatch or the visitor pattern to express operations on two dynamic types (unless you need an optimized lookup); §22.3.1.
- Don't call virtual functions during construction or destruction; §22.4.
- Use `typeid` to implement extended type information; §22.5.1.
- Use `typeid` to find the type of an object (and not to find an interface to an object); §22.5.
- Prefer virtual functions to repeated switch-statements based on `typeid` or `dynamic_cast`; §22.6

TEMPLATES

- Use templates to express algorithms that apply to many argument types; §23.1.
- Use templates to express containers; §23.2.
- Note that `template<class T>` and `template<typename T>` are synonymous; §23.2.
- When defining a template, first design and debug a non-template version; later generalize by adding parameters; §23.2.1.
- Templates are type-safe, but checking happens too late; §23.3.
- When designing a template, carefully consider the concepts (requirements) assumed for its template arguments; §23.3.
- If a class template should be copyable, give it a non-template copy constructor and a non-template copy assignment; §23.4.6.1.
- If a class template should be movable, give it a non-template move constructor and a non-template move assignment; §23.4.6.1.
- A virtual function member cannot be a template member function; §23.4.6.2.
- Define a type as a member of a template only if it depends on all the class template's arguments; §23.4.6.3.
- Use function templates to deduce class template argument types; §23.5.1.
- Overload function templates to get the same semantics for a variety of argument types; §23.5.3.
- Use argument substitution failure to provide just the right set of functions for a program; §23.5.3.2.
- Use template aliases to simplify notation and hide implementation details; §23.6.
- There is no separate compilation of templates: `#include` template definitions in every translation unit that uses them; §23.7.
- Use ordinary functions as interfaces to code that cannot deal with templates; §23.7.1.
- Separately compile large templates and templates with nontrivial context dependencies; §23.7

GENERIC PROGRAMMING

- A template can pass argument types without loss of information; §24.1.
- Templates provide a general mechanism for compile-time programming; §24.1.
- Templates provide compile-time “duck typing”; §24.1.
- Design generic algorithms by “lifting” from concrete examples; §24.2.
- Generalize algorithms by specifying template argument requirements in terms of concepts; §24.3.
- Do not give unconventional meaning to conventional notation; §24.3.
- Use concepts as a design tool; §24.3.
- Aim for “plug compatibility” among algorithms and argument type by using common and regular template argument requirements; §24.3.
- Discover a concept by minimizing an algorithm's requirements on its template arguments and then generalizing for wider use; §24.3.1.

- A concept is not just a description of the needs of a particular implementation of an algorithm; §24.3.1.
- If possible, choose a concept from a list of well-known concepts; §24.3.1, §24.4.4.
- The default concept for a template argument is Regular; §24.3.1.
- Not all template argument types are Regular; §24.3.1.
- A concept requires a semantic aspect; it is not primarily a syntactic notion; §24.3.1, §24.3.2, §24.4.1.
- Make concepts concrete in code; §24.4.
- Express concepts as compile-time predicates (constexpr functions) and test them using `static_assert()` or `enable_if<>`; §24.4.
- Use axioms as a design tool; §24.4.1.
- Use axioms as a guide for testing; §24.4.1.
- Some concepts involve two or more template arguments; §24.4.2.
- Concepts are not just types of types; §24.4.2.
- Concepts can involve numeric values; §24.4.3.
- Use concepts as a guide for testing template definitions; §24.4.5

SPECIALIZATION

- Use templates to improve type safety; §25.1.
- Use templates to raise the level of abstraction of code; §25.1.
- Use templates to provide flexible and efficient parameterization of types and algorithms; §25.1.
- Remember that value template arguments must be compile-time constants; §25.2.2.
- Use function objects as type arguments to parameterize types and algorithms with “policies”; §25.2.3.
- Use default template arguments to provide simple notation for simple uses; §25.2.5.
- Specialize templates for irregular types (such as arrays); §25.3.
- Specialize templates to optimize for important cases; §25.3.
- Define the primary template before any specialization; §25.3.1.1.
- A specialization must be in scope for every use; §25.3.1.1.

INSTANTIATION

- Let the compiler/implementation generate specializations as needed; §26.2.1.
- Explicitly instantiate if you need exact control of the instantiation environment; §26.2.2.
- Explicitly instantiate if you optimize the time needed to generate specializations; §26.2.2.
- Avoid subtle context dependencies in a template definition; §26.3.
- Names must be in scope when used in a template definition or findable through argument-dependent lookup (ADL); §26.3, §26.3.5.
- Keep the binding context unchanged between instantiation points; §26.3.4.
- Avoid fully general templates that can be found by ADL; §26.3.6.
- Use concepts and/or `static_assert` to avoid using inappropriate templates; §26.3.6.
- Use using-declarations to limit the reach of ADL; §26.3.6.
- Qualify names from a template base class `with->orT::as appropriate`; §26.3.7.

TEMPLATES AND HIERARCHIES

- When having to express a general idea in code, consider whether to represent it as a template or as a class hierarchy; §27.1.
- A template usually provides common code for a variety of arguments; §27.1.
- An abstract class can completely hide implementation details from users; §27.1.
- Irregular implementations are usually best represented as derived classes; §27.2.
- If explicit use of free store is undesirable, templates have an advantage over class hierarchies; §27.2.
- Templates have an advantage over abstract classes where inlining is important; §27.2.
- Template interfaces are easily expressed in terms of template argument types; §27.2.
- If run-time resolution is needed, class hierarchies are necessary; §27.2.
- The combination of templates and class hierarchies is often superior to either without the other; §27.2.
- Think of templates as type generators (and function generators); §27.2.1.
- There is no default relation between two classes generated from the same template; §27.2.1.
- Do not mix class hierarchies and arrays; §27.2.1.
- Do not naively templatize large class hierarchies; §27.3.
- A template can be used to provide a type-safe interface to a single (weakly typed) implementation; §27.3.1.

- Templates can be used to compose type-safe and compact data structures; §27.4.1.
- Templates can be used to linearize a class hierarchy (minimizing space and access time); §27.4.2.

METAPROGRAMMING

- Use metaprogramming to improve type safety; §28.1.
- Use metaprogramming to improve performance by moving computation to compile time; §28.1.
- Avoid using metaprogramming to an extent where it significantly slows down compilation; §28.1.
- Think in terms of compile-time evaluation and type functions; §28.2.
- Use template aliases as the interfaces to type functions returning types; §28.2.1.
- Use constexpr functions as the interfaces to type functions returning (non-type) values; §28.2.2.
- Use traits to nonintrusively associate properties with types; §28.2.4.
- Use Conditional to choose between two types; §28.3.1.1.
- Use Select to choose among several alternative types; §28.3.1.3.
- Use recursion to express compile-time iteration; §28.3.2.
- Use metaprogramming for tasks that cannot be done well at run time; §28.3.3.
- Use Enable_if to selectively declare function templates; §28.4.
- Concepts are among the most useful predicates to use with Enable_if; §28.4.3.
- Use variadic templates when you need a function that takes a variable number of arguments of a variety of types; §28.6.
- Don't use variadic templates for homogeneous argument lists (prefer initializer lists for that); §28.6.
- Use variadic templates and std::move() where forwarding is needed; §28.6.3.
- Use simple metaprogramming to implement efficient and elegant unit systems (for fine-grained type checking); §28.7.
- Use user-defined literals to simplify the use of units; §28.7

A MATRIX DESIGN

- List basic use cases; §29.1.1.
- Always provide input and output operations to simplify simple testing (e.g., unit testing); §29.1.1.
- Carefully list the properties a program, class, or library ideally should have; §29.1.2.
- List the properties of a program, class, or library that are considered beyond the scope of the project; §29.1.2.
- When designing a container template, carefully consider the requirements on the element type; §29.1.2.
- Consider how the design might accommodate run-time checking (e.g., for debugging); §29.1.2.
- If possible, design a class to mimic existing professional notation and semantics; §29.1.2.
- Make sure that the design does not leak resources (e.g., have a unique owner for each resource and use RAII); §29.2.
- Consider how a class can be constructed and copied; §29.1.1.
- Provide complete, flexible, efficient, and semantically meaningful access to elements; §29.2.2, §29.3.
- Place implementation details in their own _impl namespace; §29.4.
- Provide common operations that do not require direct access to the representation as helper functions; §29.3.2, §29.3.3.
- For fast access, keep data compact and use accessor objects to provide necessary non trivial access operations; §29.4.1, §29.4.2, §29.4.3.
- The structure of data can often be expressed as nested initializer lists; §29.4.4.
- When dealing with numbers, a way's consider "end cases," such as zero and "many"; §29.4.6.
- In addition to unit testing and testing that the code meets its requirements, test the design through examples of real use; §29.5.
- Consider how the design might accommodate unusually stringent performance requirements; §29.5.4

THE STANDARD LIBRARY

STL SUMMARY

- Use standard-library facilities to maintain portability; §30.1, §30.1.1.
- Use standard-library facilities to minimize maintenance costs; §30.1.
- Use standard-library facilities as a base for more extensive and more specialized libraries; §30.1.1.
- Use standard-library facilities as a model for flexible, widely usable software; §30.1.1.
- The standard-library facilities are defined in namespace std and found in standard-library headers; §30.2.
- A C standard-library headerX.h is presented as a C++ standard-library header in <cX>; §30.2.
- Do not try to use a standard-library facility without #includeing its header; §30.2.
- To use a range-for on a built-in array, #include<iterator>; §30.3.2.
- Prefer exception-based error handling over return-code-based error handling; §30.4.

- Always catch exception& (for standard-library and language support exceptions) and...(for unexpected exceptions); §30.4.1.
- The standard-library exception hierarchy can be (but does not have to be) used for a user's own exceptions; §30.4.1.1.
- Call terminate() in case of serious trouble; §30.4.1.3.
- Use static_assert() and assert() extensively; §30.4.2.
- Do not assume that assert() is always evaluated; §30.4.2.
- If you can't use exceptions, consider <system_error>; §30.4.3.

STL CONTAINERS

- An STL container defines a sequence; §31.2.
- Use vector as your default container; §31.2, §31.4.
- Insertion operators, such as insert() and push_back() are often more efficient on a vector than on a list; §31.2, §31.4.1.1.
- Use forward_list for sequences that are usually empty; §31.2, §31.4.2.
- When it comes to performance, don't trust your intuition: measure; §31.3.
- Don't blindly trust asymptotic complexity measures; some sequences are short and the cost of individual operations can vary dramatically; §31.3.
- STL containers are resource handles; §31.2.1.
- A map is usually implemented as a red-black tree; §31.2.1, §31.4.3.
- An unordered_map is a hash table; §31.2.1, §31.4.3.2.
- To be an element type for a STL container, a type must provide copy or move operations; §31.2.2.
- Use containers of pointers or smart pointers when you need to preserve polymorphic behavior; §31.2.2.
- Comparison operations should implement a strict weak order; §31.2.2.1.
- Pass a container by reference and return a container by value; §31.3.2.
- For a container, use the()-initializer syntax for sizes and the{}-initializer syntax for lists of elements; §31.3.2.
- For simple traversals of a container, use a range-for loop or a begin/end pair of iterators; §31.3.4.
- Use const iterators where you don't need to modify the elements of a container; §31.3.4.
- Use auto to avoid verbosity and typos when you use iterators; §31.3.4.
- Use reserve() to avoid invalidating pointers and iterators to elements; §31.3.3, §31.4.1.
- Don't assume performance benefits from reserve() without measurement; §31.3.3.
- Use push_back() or resize() on a container rather than realloc() on an array; §31.3.3, §31.4.1.1.
- Don't use iterators into a resized vector or deque; §31.3.3.
- When necessary, use reserve() to make performance predictable; §31.3.3.
- Do not assume that [] range checks; §31.2.2.
- Use at() when you need guaranteed range checks; §31.2.2.
- Use emplace() for notational convenience; §31.3.7

STL ALGORITHMS

- An STL algorithm operates on one or more sequences; §32.2.
- An input sequence is half-open and defined by a pair of iterators; §32.2.
- When searching, an algorithm usually returns the end of the input sequence to indicate "not found"; §32.2.
- Prefer a carefully specified algorithm to "random code"; §32.2.
- When writing a loop, consider whether it could be expressed as a general algorithm; §32.2.
- Make sure that a pair of iterator arguments really do specify a sequence; §32.2.
- When the pair-of-iterators style becomes tedious, introduce a container/range algorithm; §32.2.
- Use predicates and other function objects to give standard algorithms a wider range of meanings; §32.3.
- A predicate must not modify its argument; §32.3.
- The default == and < on pointers are rarely adequate for standard algorithms; §32.3.
- Know the complexity of the algorithms you use, but remember that a complexity measure is only a rough guide to performance; §32.3.1.
- Use for_each() and transform() only when there is no more-specific algorithm for a task; §32.4.1.
- Algorithms do not directly add or subtract elements from their argument sequences; §32.5.2, §32.5.3.
- If you have to deal with uninitialized objects, consider the uninitialized_* algorithms; §32.5.6.
- An STL algorithm uses an equality comparison generated from its ordering comparison, rather than==; §32.6.
- Note that sorting and searching C-style strings requires the user to supply a string comparison operation; §32.6

STL ITERATORS

- An input sequence is defined by a pair of iterators; §33.1.1.

- An output sequence is defined by a single iterator; avoid overflow; §33.1.1.
- For any iterator `p`, `[p:p]` is the empty sequence; §33.1.1.
- Use the end of a sequence to indicate “not found”; §33.1.1.
- Think of iterators as more general and often better behaved pointers; §33.1.1.
- Use iterator types, such as `list<char>::iterator`, rather than pointers to refer to elements of a container; §33.1.1.
- Use `iterator_traits` to obtain information about iterators; §33.1.3.
- You can do compile-time dispatch using `iterator_traits`; §33.1.3.
- Use `iterator_traits` to select an optimal algorithm based on an iterator’s category; §33.1.3.
- `iterator_traits` are an implementation detail; prefer to use them implicitly; §33.1.3.
- Use `base()` to extract an iterator from a `reverse_iterator`; §33.2.1.
- You can use an insert iterator to add elements to a container; §33.2.2.
- A `move_iterator` can be used to make copy operations into move operations; §33.2.3.
- Make sure that your containers can be traversed using a range-for; §33.3.
- Use `bind()` to create variants of functions and function objects; §33.5.1.
- Note that `bind()` dereferences references early; use `ref()` if you want to delay dereferencing; §33.5.1.
- A `mem_fn()` or a lambda can be used to convert the `p->f(a)` calling convention into `f(p,a)`; §33.5.2.
- Use function when you need a variable that can hold a variety of callable objects; §33.5.3

MEMORY AND RESOURCES

- Use array where you need a sequence with a `constexpr` size; §34.2.1.
- Prefer array over built-in arrays; §34.2.1.
- Use `bitset` if you need `N` bits and `N` is not necessarily the number of bits in a built-in integer type; §34.2.2.
- Avoid `vector<bool>`; §34.2.3.
- When using `pair`, consider `make_pair()` for type deduction; §34.2.4.1.
- When using `tuple`, consider `make_tuple()` for type deduction; §34.2.4.2.
- Use `unique_ptr` to represent exclusive ownership; §34.3.1.
- Use `shared_ptr` to represent shared ownership; §34.3.2.
- Minimize the use of `weak_ptrs`; §34.3.3.
- Use allocators (only) when the usual `new/delete` semantics is insufficient for logical or performance reasons; §34.4.
- Prefer resource handles with specific semantics to smart pointers; §34.5.
- Prefer `unique_ptr` to `shared_ptr`; §34.5.
- Prefer smart pointers to garbage collection; §34.5.
- Have a coherent and complete strategy for management of general resources; §34.5.
- Garbage collection can be really useful for dealing with leaks in programs with messy pointer use; §34.5.
- Garbage collection is optional; §34.5.
- Don’t disguise pointers (even if you don’t use garbage collection); §34.5.
- If you use garbage collection, use `declare_no_pointers()` to let the garbage collector ignore data that cannot contain pointers; §34.5.
- Don’t mess with uninitialized memory unless you absolutely have to; §34.6

UTILITIES

- Use `<chrono>` facilities, such as `steady_clock`, `duration`, and `time_point` for timing; §35.2.
- Prefer `<clock>` facilities over `<ctime>` facilities; §35.2.
- Use `duration_cast` to get durations in known units of time; §35.2.1.
- Use `system_clock::now()` to get the current time; §35.2.3.
- You can inquire about properties of types at compile time; §35.4.1.
- Use `move(obj)` only when the value of `obj` cannot be used again; §35.5.1.
- Use `forward()` for forwarding; §35.5.1

STRINGS

- Use character classifications rather than handcrafted checks on character ranges; §36.2.1.
- If you implement string-like abstractions, use `character_traits` to implement operations on characters; §36.2.2.
- A `basic_string` can be used to make strings of characters on any type; §36.3.
- Use strings as variables and members rather than as base classes; §36.3.
- Prefer string operations to C-style string functions; §36.3.1.
- Return strings by value (rely on move semantics); §36.3.2.

- Use `string::npos` to indicate “the rest of the string”; §36.3.2.
- Do not pass a `nullptr` to a string function expecting a C-style string; §36.3.2.
- A string can grow and shrink, as needed; §36.3.3.
- Use `at()` rather than iterators or `[]` when you want range checking; §36.3.3, §36.3.6.
- Use iterators and `[]` rather than `at()` when you want to optimize speed; §36.3.3, §36.3.6.
- If you use strings, catch `length_error` and `out_of_range` somewhere; §36.3.3.
- Use `c_str()` to produce a C-style string representation of a string (only) when you have to; §36.3.3.
- string input is type sensitive and doesn't overflow; §36.3.4.
- Prefer a `string_stream` or a generic value extraction function (such as `to<X>()`) over direct use of `str*` numeric conversion functions; §36.3.5.
- Use the `find()` operations to locate values in a string (rather than writing an explicit loop); §36.3.7.
- Directly or indirectly, use `substr()` to read substrings and `replace()` to write substrings; §36.3.8

REGULAR EXPRESSIONS

- Use `regex` for most conventional uses of regular expressions; §37.1.
- The regular expression notation can be adjusted to match various standards; §37.1.1, §37.2.
- The default regular expression notation is that of ECMAScript; §37.1.1.
- For portability, use the character class notation to avoid nonstandard abbreviations; §37.1.1.
- Be restrained; regular expressions can easily become a write-only language; §37.1.1.
- Prefer raw string literals for expressing all but the simplest patterns; §37.1.1.
- Note that `\i` allows you to express a subpattern in terms of a previous subpattern; §37.1.1.
- Use `?` to make patterns “lazy”; §37.1.1, §37.2.1.
- `regex` can use ECMAScript, POSIX, `awk`, `grep`, and `egrep` notation; §37.2.
- Keep a copy of the pattern string in case you need to output it; §37.2.
- Use `regex_search()` for looking at streams of characters and `regex_match()` to look for fixed layouts; §37.3.2, §37.3.1

I/O STREAMS

- Define `<<` and `>>` or user-defined types with values that have meaningful textual representations; §38.1, §38.4.1, §38.4.2.
- Use `cout` for normal output and `cerr` for errors; §38.1.
- There are `iostreams` for ordinary characters and wide characters, and you can define an `iostream` for any kind of character; §38.1.
- There are standard `iostreams` for standard I/O streams, files, and strings; §38.2.
- Don't try to copy a file stream; §38.2.1.
- Binary I/O is system specific; §38.2.1.
- Remember to check that a file stream is attached to a file before using it; §38.2.1.
- Prefer `ifstream` and `ofstream` over the generic `fstream`; §38.2.1.
- Use `stringstreams` for in-memory formatting; §38.2.2.
- Use exceptions to catch rare `bad()` I/O errors; §38.3.
- Use the stream state `fail` to handle potentially recoverable I/O errors; §38.3.
- You don't need to modify `istream` or `ostream` to add new `<<` and `>>` operators; §38.4.1.
- When implementing a `iostream` primitive operation, use `sentry`; §38.4.1.
- Prefer formatted input over unformatted, low-level input; §38.4.1.
- Input into strings does not overflow; §38.4.1.
- Be careful with the termination criteria when using `get()`, `getline()`, and `read()`; §38.4.1.
- By default `>>` skips whitespace; §38.4.1.
- You can define a `<<` (or a `>>`) so that it behaves as a virtual function based on its second operand; §38.4.2.1.
- Prefer manipulators to state flags for controlling I/O; §38.4.3.
- Use `sync_with_stdio(true)` if you want to mix C-style and `iostream` I/O; §38.4.4.
- Use `sync_with_stdio(false)` to optimize `iostreams`; §38.4.4.
- Tie streams used for interactive I/O; §38.4.4.
- Use `imbue()` to make an `iostream` reflect “cultural differences” of a locale; §38.4.4.
- `width()` specifications apply to the immediately following I/O operation only; §38.4.5.1.
- `precision()` specifications apply to all following floating-point output operations; §38.4.5.1.
- Floating-point format specifications (e.g., `scientific`) apply to all following floating-point out-put operations; §38.4.5.2.
- `#include <iomanip>` when using standard manipulators taking arguments; §38.4.5.2.
- You hardly ever need to flush(); §38.4.5.2.

- Don't use endl except possibly for aesthetic reasons; §38.4.5.2.
- If iostream formatting gets too tedious, write your own manipulators; §38.4.5.3.
- You can achieve the effect (and efficiency) of a ternary operator by defining a simple function object; §38.4.5.3

LOCALES

- Expect that every nontrivial program or system that interacts directly with people will be used in several different countries; §39.1.
- Don't assume that everyone uses the same character set as you do; §39.1, §39.4.1.
- Prefer using locales to writing ad hoc code for culture-sensitive I/O; §39.1.
- Use locales to meet external (non-C++) standards; §39.1.
- Think of a locale as a container of facets; §39.2.
- Avoid embedding locale name strings in program text; §39.2.1.
- Keep changes of locale to a few places in a program; §39.2.1.
- Minimize the use of global format information; §39.2.1.
- Prefer locale-sensitive string comparisons and sorts; §39.2.2, §39.4.1.
- Make facets immutable; §39.3.
- Let locale handle the lifetime of facets; §39.3.
- You can make your own facets; §39.3.2.
- When writing locale-sensitive I/O functions, remember to handle exceptions from user-supplied (overriding) functions; §39.4.2.2.
- Use numput if you need separators in numbers; §39.4.2.1.
- Use a simple Money type to hold monetary values; §39.4.3.
- Use simple user-defined types to hold values that require locale-sensitive I/O (rather than casting to and from values of built-in types); §39.4.3.
- The time_put facet can be used for both <chrono>- and <ctime>-style time §39.4.4.
- Prefer the character classification functions in which the locale is explicit; §39.4.5, §39.5.

NUMERICS

- Numerical problems are often subtle. If you are not 100% certain about the mathematical aspects of a numerical problem, either take expert advice, experiment, or do both; §29.1.
- Use variants of numeric types that are appropriate for their use; §40.2.
- Use numeric_limits to check that the numeric types are adequate for their use; §40.2.
- Specialize numeric_limits for a user-defined numeric type; §40.2.
- Prefer numeric_limits over limit macros; §40.2.1.
- Use std::complex for complex arithmetic; §40.4.
- Use {}-initialization to protect against narrowing; §40.4.
- Use valarray for numeric computation when run-time efficiency is more important than flexibility with respect to operations and element types; §40.5.
- Express operations on part of an array in terms of slices rather than loops; §40.5.5.
- Slices is a generally useful abstraction for access of compact data; §40.5.4, §40.5.6.
- Consider accumulate(), inner_product(), partial_sum(), and adjacent_difference() before you write a loop to compute a value from a sequence; §40.6.
- Bind an engine to a distribution to get a random number generator; §40.7.
- Be careful that your random numbers are sufficiently random; §40.7.1.
- If you need genuinely random numbers (not just a pseudo-random sequence), use random_device; §40.7.2.
- Prefer a random number class for a particular distribution over direct use of rand(); §40.7.4.

CONCURRENCY

- Use concurrency to improve responsiveness or to improve throughput; §41.1.
- Work at the highest level of abstraction that you can afford; §41.1.
- Prefer packaged_task and futures over direct use of threads and mutexes; §41.1.
- Prefer mutexes and condition_variables over direct use of atomics except for simple counters; §41.1.
- Avoid explicitly shared data whenever you can; §41.1.
- Consider processes as an alternative to threads; §41.1.
- The standard-library concurrency facilities are type safe; §41.1.
- The memory model exists to save most programmers from having to think about the machine architecture level of computers; §41.2.

- The memory model makes memory appear roughly as naively expected; §41.2.
- Separate threads accessing separate bit-fields of a struct may interfere with each other; §41.2.
- Avoid data races; §41.2.4.
- Atomics allow for lock-free programming; §41.3.
- Lock-free programming can be essential for avoiding deadlock and to ensure that every thread makes progress; §41.3.
- Leave lock-free programming to experts; §41.3.
- Leave relaxed memory models to experts; §41.3.
- A volatile tells the compiler that the value of an object can be changed by something that is not part of the program; §41.4.
- A C++ volatile is not a synchronization mechanism; §41.4.

THREADS AND TASKS

- A thread is a type-safe interface to a system thread; §42.2.
- Do not destroy a running thread; §42.2.2.
- Use join() to wait for a thread to complete; §42.2.4.
- Consider using a guarded_thread to provide RAI for threads; §42.2.4.
- Do not detach() a thread unless you absolutely have to; §42.2.4.
- Use lock_guard or unique_lock to manage mutexes; §42.3.1.4.
- Use lock() to acquire multiple locks; §42.3.2.
- Use condition_variables to manage communication among threads; §42.3.4.
- Think in terms of tasks that can be executed concurrently, rather than directly in terms of threads; §42.4.
- Value simplicity; §42.4.
- Return a result using a promise and get a result from a future; §42.4.1.
- Don't set_value() or set_exception() to a promise twice; §42.4.2.
- Use packaged_tasks to handle exceptions thrown by tasks and to arrange for value return; §42.4.3.
- Use a packaged_task and a future to express a request to an external service and wait for its response; §42.4.3.
- Don't get() twice from a future; §42.4.4.
- Use async() to launch simple tasks; §42.4.6.
- Picking a good granularity of concurrent tasks is difficult: experiment and measure; §42.4.7.
- Whenever possible, hide concurrency behind the interface of a parallel algorithm; §42.4.7.
- A parallel algorithm may be semantically different from a sequential solution to the same problem (e.g., pfind_all() vs. find()); §42.4.7.
- Sometimes, a sequential solution is simpler and faster than a concurrent solution; §42.4.7.

THE C STANDARD LIBRARY

- Use fstreams rather than fopen()/fclose() if you worry about resource leaks; §43.2.
- Prefer <iostream> to <stdlib> for reasons of type safety and extensibility; §43.3.
- Never use gets() or scanf("%s",s); §43.3.
- Prefer <string> to <cstring> for reasons of ease of use and simplicity of resource management; §43.4.
- Use the C memory management routines, such as memcpy(), only for raw memory; §43.5.
- Prefer vector to uses of malloc() and realloc(); §43.5.
- Beware that the C standard library does not know about constructors and destructors; §43.5.
- Prefer <chrono> to <ctime> for timing; §43.6.
- For flexibility, ease of use, and performance, prefer sort() over qsort(); §43.7.
- Don't use exit(); instead, throw an exception; §43.7.
- Don't use longjmp(); instead, throw an exception; §43.7.

COMPATIBILITY

- Before using a new feature in production code, try it out by writing small programs to test the standards conformance and performance of the implementations you plan to use; §44.1.
- For learning C++, use the most up-to-date and complete implementation of Standard C++ that you can get access to; §44.2.4.
- The common subset of C and C++ is not the best initial subset of C++ to learn; §1.2.3, §44.2.4.
- Prefer standard facilities to nonstandard ones; §36.1, §44.2.4.
- Avoid deprecated features such as throw-specifications; §44.2.3, §13.5.1.3.
- Avoid C-style casts; §44.2.3, §11.5. "Implicit int" has been banned, so explicitly specify the type of every function, variable, const, etc.; §44.3.3.

- When converting a C program to C++, first make sure that function declarations (prototypes) and standard headers are used consistently; §44.3.3.
- When converting a C program to C++, rename variables that are C++ keywords; §44.3.3.
- For portability and type safety, if you must use C, write in the common subset of C and C++; §44.2.4.
- When converting a C program to C++, cast the result of malloc() to the proper type or change all uses of malloc() to uses of new; §44.3.3.
- When converting from malloc() and free() to new and delete, consider using vector, push_back(), and reserve() instead of realloc(); §3.4.2, §43.5.
- When converting a C program to C++, remember that there are no implicit conversions from ints to enumerations; use explicit type conversion where necessary; §44.3.3, §8.4.
- A facility defined in namespace std is defined in a header without a suffix (e.g. std::cout is declared in <iostream>); §30.2.
- Use <string> to get std::string (<string.h> holds the C-style string functions); §15.2.4.
- For each standard C header <X.h> that places names in the global namespace, the header <cX> places the names in namespace std; §15.2.2.
- Use extern "C" when declaring C functions; §15.2.5